

1. Import Libraries

```
# Import necessary libraries for data manipulation, visualization, and machine learning.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import (
    accuracy_score, classification_report, roc_auc_score,
    roc_curve, confusion_matrix
)
from imblearn.over_sampling import SMOTE
```

2. Load and Inspect Data

```
# Load the dataset
df = pd.read_csv('creditcard.csv')

df.head()
```

	Time	V1	V2	V3	V4	V5	V6
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921

	V8	V9	...	V21	V22	V23	V24
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575

```
0.647376
4 -0.270533  0.817739  ... -0.009431  0.798278 -0.137458  0.141267 -
0.206010
```

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

```
[5 rows x 31 columns]
```

```
df.shape
```

```
(284807, 31)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	Time	284807 non-null	float64
1	V1	284807 non-null	float64
2	V2	284807 non-null	float64
3	V3	284807 non-null	float64
4	V4	284807 non-null	float64
5	V5	284807 non-null	float64
6	V6	284807 non-null	float64
7	V7	284807 non-null	float64
8	V8	284807 non-null	float64
9	V9	284807 non-null	float64
10	V10	284807 non-null	float64
11	V11	284807 non-null	float64
12	V12	284807 non-null	float64
13	V13	284807 non-null	float64
14	V14	284807 non-null	float64
15	V15	284807 non-null	float64
16	V16	284807 non-null	float64
17	V17	284807 non-null	float64
18	V18	284807 non-null	float64
19	V19	284807 non-null	float64
20	V20	284807 non-null	float64
21	V21	284807 non-null	float64
22	V22	284807 non-null	float64
23	V23	284807 non-null	float64
24	V24	284807 non-null	float64
25	V25	284807 non-null	float64

```
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount   284807 non-null float64
30 Class    284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
# Check for missing values
df.isna().sum()
```

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
```

3. Preprocess Data

```
# Standardize the 'Amount' column, drop the 'Time' column, and handle
duplicates.
```

```
sc = StandardScaler()
df[ 'Amount'] = sc. fit_transform(pd. DataFrame (df[ 'Amount']))
df.head()
```

	Time	V1	V2	V3	V4	V5	V6
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921

	V8	V9	...	V21	V22	V23	V24
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	0.244964	0
1	0.125895	-0.008983	0.014724	-0.342475	0
2	-0.139097	-0.055353	-0.059752	1.160686	0
3	-0.221929	0.062723	0.061458	0.140534	0
4	0.502292	0.219422	0.215153	-0.073403	0

```
[5 rows x 31 columns]
```

```
# Drop the 'Time' column as it is not relevant for this analysis
```

```
df = df.drop(['Time'], axis=1)
```

```
# Check for duplicates
```

```
df.duplicated().any()
```

```
np.True_
```

```
print(df.shape)
```

```
df = df.drop_duplicates()
```

```
print(df.shape)
```

```

(284807, 30)
(275663, 30)

# Check the proportion of fraud (Class=1) vs. non-fraud (Class=0)
df['Class'].value_counts(normalize=True)

Class
0    0.998284
1    0.001716
Name: proportion, dtype: float64

```

There is a

```
df.describe().T
```

	count	mean	std	min	25%	50%
V1	275663.0	-0.037460	1.952522	-56.407510	-0.941105	-0.059659
V2	275663.0	-0.002430	1.667260	-72.715728	-0.614040	0.070249
V3	275663.0	0.025520	1.507538	-48.325589	-0.843168	0.200736
V4	275663.0	-0.004359	1.424323	-5.683171	-0.862847	-0.035098
V5	275663.0	-0.010660	1.378117	-113.743307	-0.700192	-0.060556
V6	275663.0	-0.014206	1.313213	-26.160506	-0.765861	-0.270931
V7	275663.0	0.008586	1.240348	-43.557242	-0.552047	0.044848
V8	275663.0	-0.005698	1.191596	-73.216718	-0.209618	0.022980
V9	275663.0	-0.012363	1.100108	-13.434066	-0.659904	-0.064724
V10	275663.0	0.003114	1.087025	-24.588262	-0.538968	-0.091752
V11	275663.0	-0.007174	1.020571	-4.797473	-0.772693	-0.039469
V12	275663.0	-0.005347	0.998661	-18.683715	-0.413717	0.133349
V13	275663.0	0.000539	0.999660	-5.791881	-0.654360	-0.011557
V14	275663.0	0.000681	0.952571	-19.214325	-0.425932	0.049552
V15	275663.0	-0.010315	0.917772	-4.498945	-0.596079	0.036145
V16	275663.0	-0.004319	0.880320	-14.129855	-0.477642	0.061670
V17	275663.0	0.000479	0.844821	-25.162799	-0.482600	-0.063489

V18	275663.0	0.003874	0.841638	-9.498746	-0.498396	0.001392
V19	275663.0	0.000511	0.820520	-7.213527	-0.464409	0.001119
V20	275663.0	0.003407	0.779950	-54.497720	-0.212027	-0.058053
V21	275663.0	0.002579	0.733089	-34.830382	-0.225021	-0.025637
V22	275663.0	0.005827	0.726378	-10.933144	-0.532173	0.013397
V23	275663.0	-0.001941	0.631451	-44.807735	-0.165440	-0.013655
V24	275663.0	-0.006868	0.605550	-2.836627	-0.361062	0.037569
V25	275663.0	-0.004812	0.524175	-10.295397	-0.323597	0.009909
V26	275663.0	-0.000240	0.484139	-2.604551	-0.328290	-0.056667
V27	275663.0	0.001921	0.401271	-22.565679	-0.071729	0.002615
V28	275663.0	0.000904	0.332649	-15.430084	-0.052654	0.011788
Amount	275663.0	0.008911	1.012371	-0.353229	-0.328041	-0.258315
Class	275663.0	0.001716	0.041388	0.000000	0.000000	0.000000
	75%	max				
V1	1.294471	2.454930				
V2	0.819067	22.057729				
V3	1.048461	9.382558				
V4	0.753943	16.875344				
V5	0.604521	34.801666				
V6	0.387704	73.301626				
V7	0.583885	120.589494				
V8	0.322319	20.007208				
V9	0.593098	15.594995				
V10	0.470702	23.745136				
V11	0.734969	12.018913				
V12	0.614670	7.848392				
V13	0.668570	7.126883				
V14	0.492169	10.526766				
V15	0.638997	8.877742				
V16	0.524709	17.315112				
V17	0.401407	9.253526				
V18	0.507708	5.041069				
V19	0.465782	5.591971				
V20	0.139803	39.420904				
V21	0.189118	27.202839				
V22	0.534272	10.503090				
V23	0.145482	22.528412				

```
V24      0.432931    4.584549
V25      0.347151    7.519589
V26      0.244196    3.517346
V27      0.094730   31.612198
V28      0.081355   33.847808
Amount -0.033742  102.362243
Class    0.000000    1.000000
```


```
sns.countplot(x='Class', data=df, palette='Set2')
```

```
/var/folders/k2/4nq8rpf91q527bjh5tbnw2v00000gn/T/
ipykernel_32475/2395609850.py:1: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Class', data=df, palette='Set2')
```

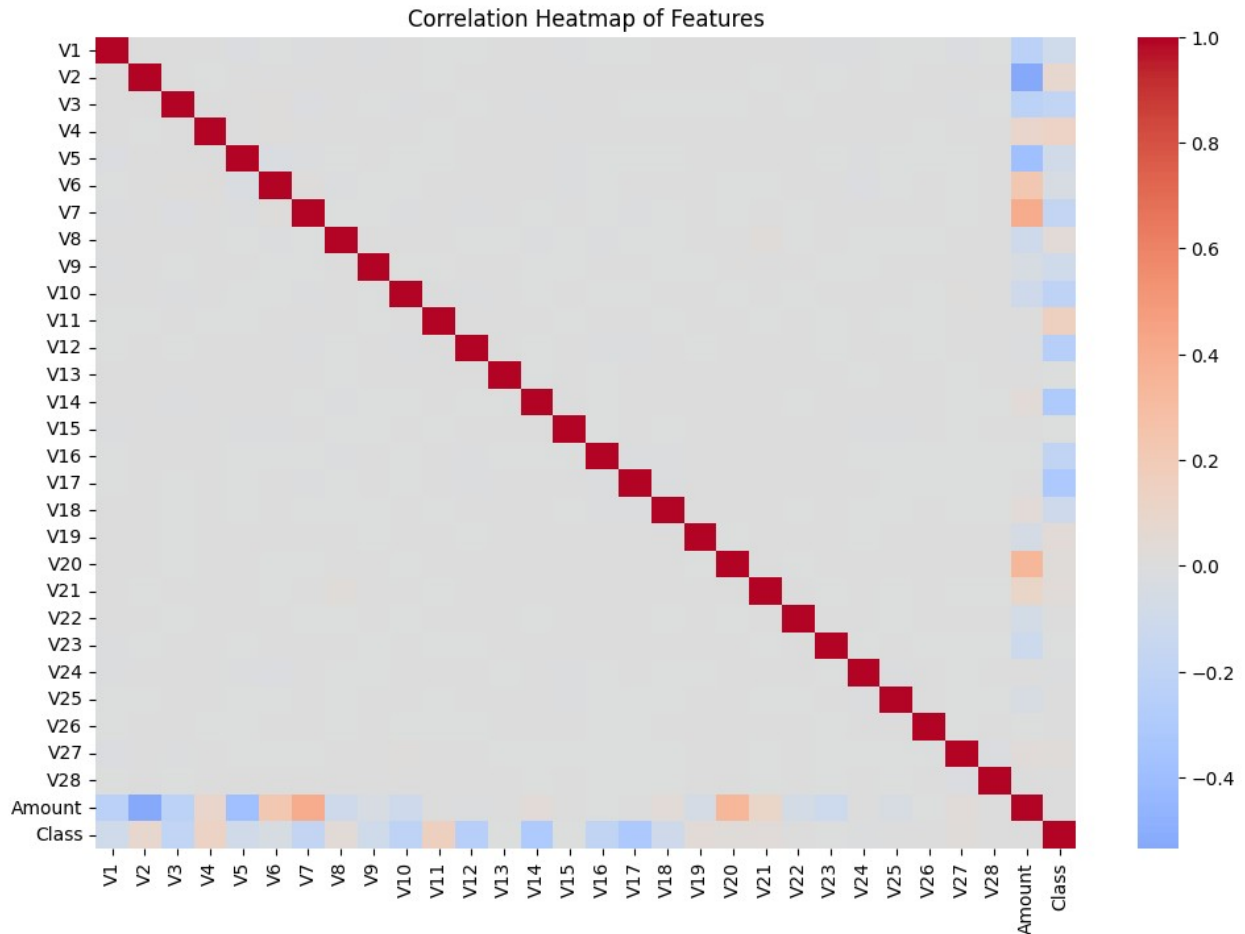
```
<Axes: xlabel='Class', ylabel='count'>
```



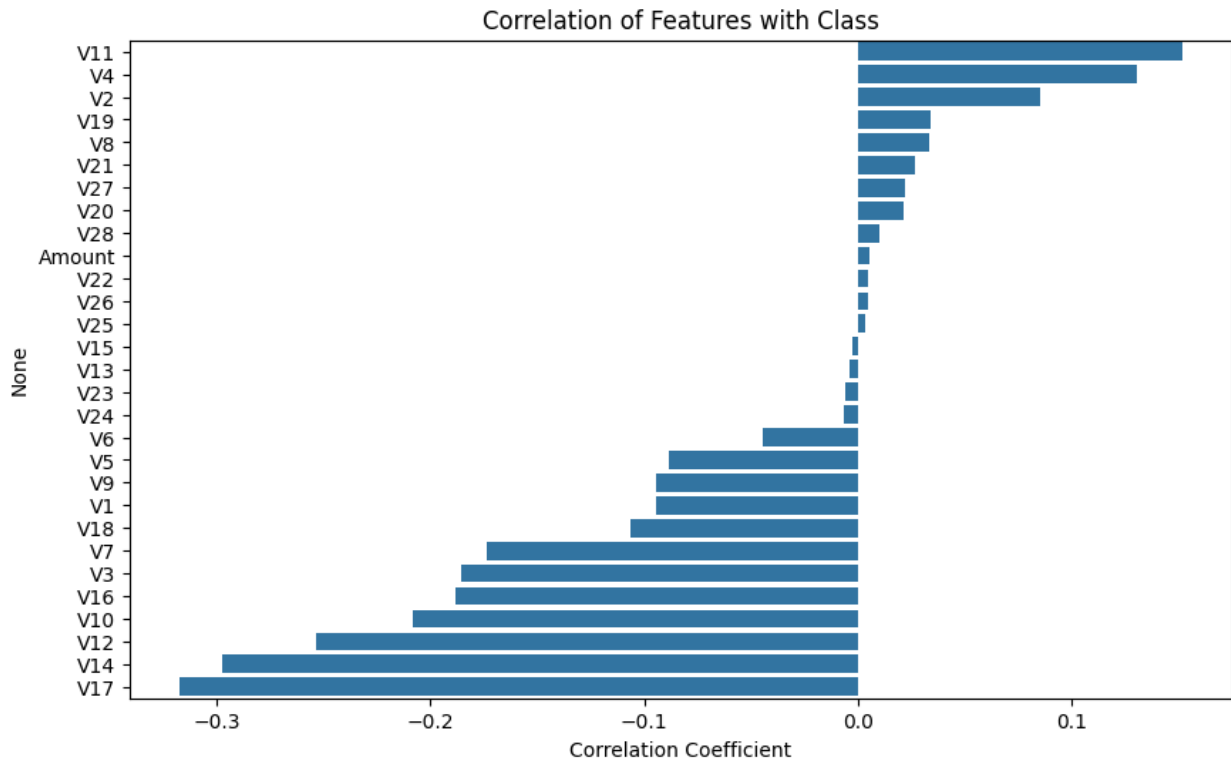
```
corr_matrix = df.corr()
```

```
plt.figure(figsize=(12, 8))
```

```
sns.heatmap(corr_matrix, cmap='coolwarm', center=0, annot=False,
fmt='.2f')
plt.title('Correlation Heatmap of Features')
plt.show()
```



```
plt.figure(figsize=(10, 6))
class_corr = corr_matrix['Class'].sort_values(ascending=False)[1:] #
Exclude Class itself
sns.barplot(x=class_corr.values, y=class_corr.index)
plt.title('Correlation of Features with Class')
plt.xlabel('Correlation Coefficient')
plt.show()
```

```
X = df.drop(['Class'], axis=1)
y = df['Class']
```

5. PCA Visualization


```
# Apply PCA
pca = PCA(n_components=2, random_state=42)
X_pca = pca.fit_transform(X)

# Create DataFrame for visualization
pca_df = pd.DataFrame({
    'PC1': X_pca[:, 0],
    'PC2': X_pca[:, 1],
    'Class': y
})

# Optional: map class labels for clearer legend
pca_df['Class'] = pca_df['Class'].map({0: 'Non-Fraud', 1: 'Fraud'})

# Plot
plt.figure(figsize=(10, 6))
sns.scatterplot(x='PC1', y='PC2', hue='Class', style='Class',
               data=pca_df, alpha=0.6)
plt.title('PCA Visualization of Transactions')
plt.xlabel(f'Principal Component 1\n({pca.explained_variance_ratio_[0]:.2%} variance)')
```

```
plt.ylabel(f'Principal Component 2  
({pca.explained_variance_ratio_[1]:.2%} variance)')
plt.legend(title='Class')
plt.show()
```



Fraud vs Non-Fraud: Fraudulent transactions (orange) form a separate cluster in PCA space, suggesting they are distinguishable from non-fraud ones.

6. Data Splitting ##### Split the data into training, validation, and test sets with stratification.

```
# Split data into 70% training and 30% test sets

X = df.drop('Class', axis=1)
y = df['Class']

X_train_raw, X_test, y_train_raw, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

# Further split training data into 80% train and 20% validation
X_train, X_val, y_train, y_val = train_test_split(
    X_train_raw, y_train_raw, test_size=0.2, random_state=42,
    stratify=y_train_raw
)
```

7. Feature Selection

```
# Train Random Forest on training data to compute feature importances
importances_model = RandomForestClassifier(random_state=42)
importances_model.fit(X_train, y_train)

feature_importances = importances_model.feature_importances_
feature_names = X_train.columns

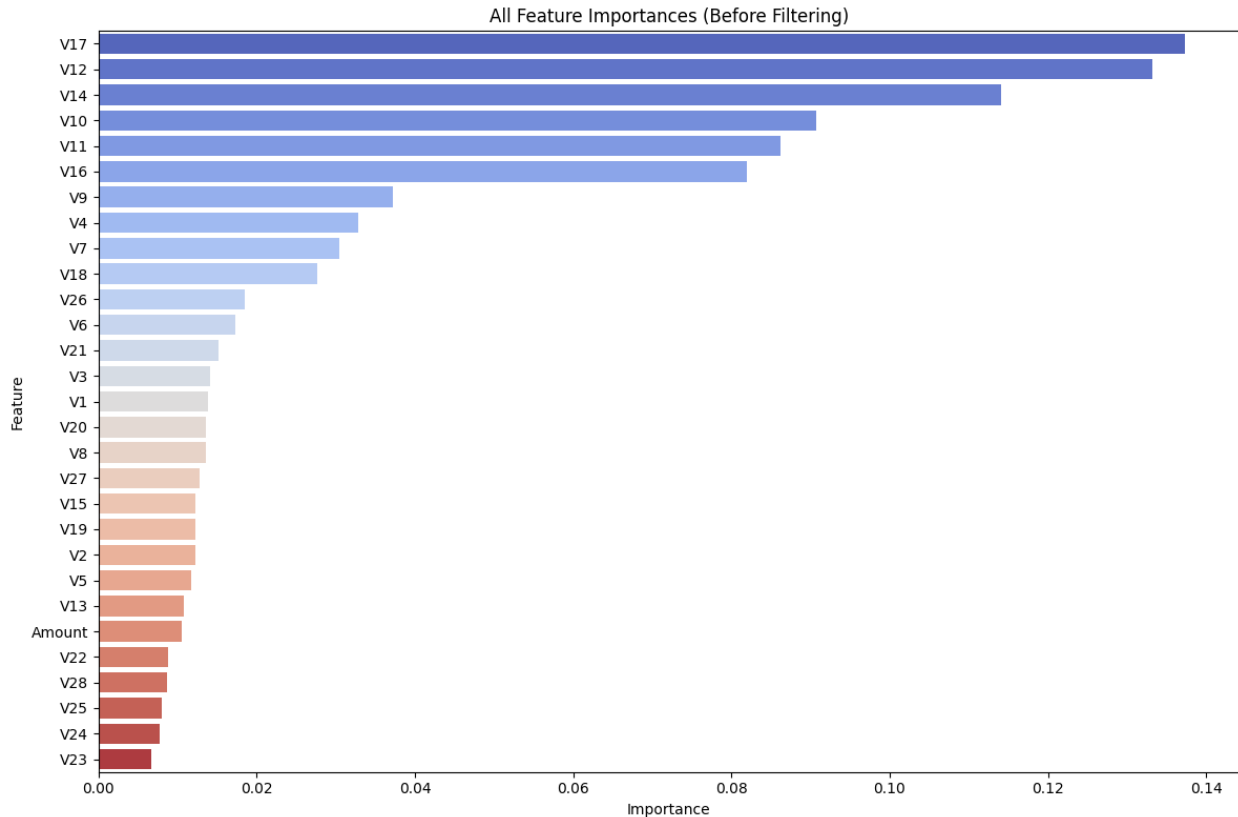
# Create DataFrame of feature importances
feature_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': feature_importances
}).sort_values(by='Importance', ascending=False)

# Visualize all feature importances
plt.figure(figsize=(12, 8))
sns.barplot(x='Importance', y='Feature', data=feature_df,
            palette='coolwarm')
plt.title("All Feature Importances (Before Filtering)")
plt.tight_layout()
plt.show()

/var/folders/k2/4nq8rpf91q527bjh5tbnw2v00000gn/T/
ipykernel_32475/3810627909.py:20: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `y` variable to `hue` and set
`legend=False` for the same effect.

    sns.barplot(x='Importance', y='Feature', data=feature_df,
    palette='coolwarm')
```



```
# Select features with importance >= 0.02
important_features = feature_df[feature_df['Importance'] >= 0.02]
['Feature']

X_train_selected = X_train[important_features]
X_val_selected = X_val[important_features]
X_test_selected = X_test[important_features]
```

8. Model Training and Evaluation (Before SMOTE)

Train and evaluate models on selected features without handling class imbalance.

```
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000,
random_state=42),
    'Decision Tree': DecisionTreeClassifier(random_state=42),
    'Random Forest': RandomForestClassifier(n_estimators=100,
random_state=42)
}

plt.figure(figsize=(10, 6))

for name, model in models.items():
    print(f"\n=== {name} ===")
```

```

model.fit(X_train_selected, y_train)

y_val_pred = model.predict(X_val_selected)
y_val_prob = model.predict_proba(X_val_selected)[: , 1]

print("--- Validation Metrics ---")
print("Val Accuracy:", accuracy_score(y_val, y_val_pred))
print("Val AUC:", roc_auc_score(y_val, y_val_prob))

y_test_pred = model.predict(X_test_selected)
y_test_prob = model.predict_proba(X_test_selected)[: , 1]

print("--- Test Metrics ---")
print("Test Accuracy:", accuracy_score(y_test, y_test_pred))
print("Classification Report:")
print(classification_report(y_test, y_test_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_test_pred))

auc = roc_auc_score(y_test, y_test_prob)
print("Test AUC Score:", auc)

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_test_prob)
plt.plot(fpr, tpr, label=f"{name} (AUC = {auc:.2f})")

# Final ROC Curve Plot
plt.plot([0, 1], [0, 1], 'k--') # Diagonal line (chance)
plt.title("ROC Curve on Test Data (Reduced Features Only)")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(True)
plt.tight_layout()
plt.show()

```

=== Logistic Regression ===

--- Validation Metrics ---

Val Accuracy: 0.9990153654807866

Val AUC: 0.9741936194294282

--- Test Metrics ---

Test Accuracy: 0.9990447284731375

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	82557
1	0.85	0.54	0.66	142
accuracy			1.00	82699

macro avg	0.93	0.77	0.83	82699
weighted avg	1.00	1.00	1.00	82699

Confusion Matrix:

```
[[82544  13]
 [   66  76]]
```

Test AUC Score: 0.9621888214834753

=== Decision Tree ===

--- Validation Metrics ---

Val Accuracy: 0.9990671883502189

Val AUC: 0.8785283205559894

--- Test Metrics ---

Test Accuracy: 0.9990689125624251

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	82557
1	0.74	0.70	0.72	142

accuracy			1.00	82699
macro avg	0.87	0.85	0.86	82699
weighted avg	1.00	1.00	1.00	82699

Confusion Matrix:

```
[[82523  34]
 [   43  99]]
```

Test AUC Score: 0.8483856309605637

=== Random Forest ===

--- Validation Metrics ---

Val Accuracy: 0.9994299484362449

Val AUC: 0.9074301296768658

--- Test Metrics ---

Test Accuracy: 0.9994195818570962

Classification Report:

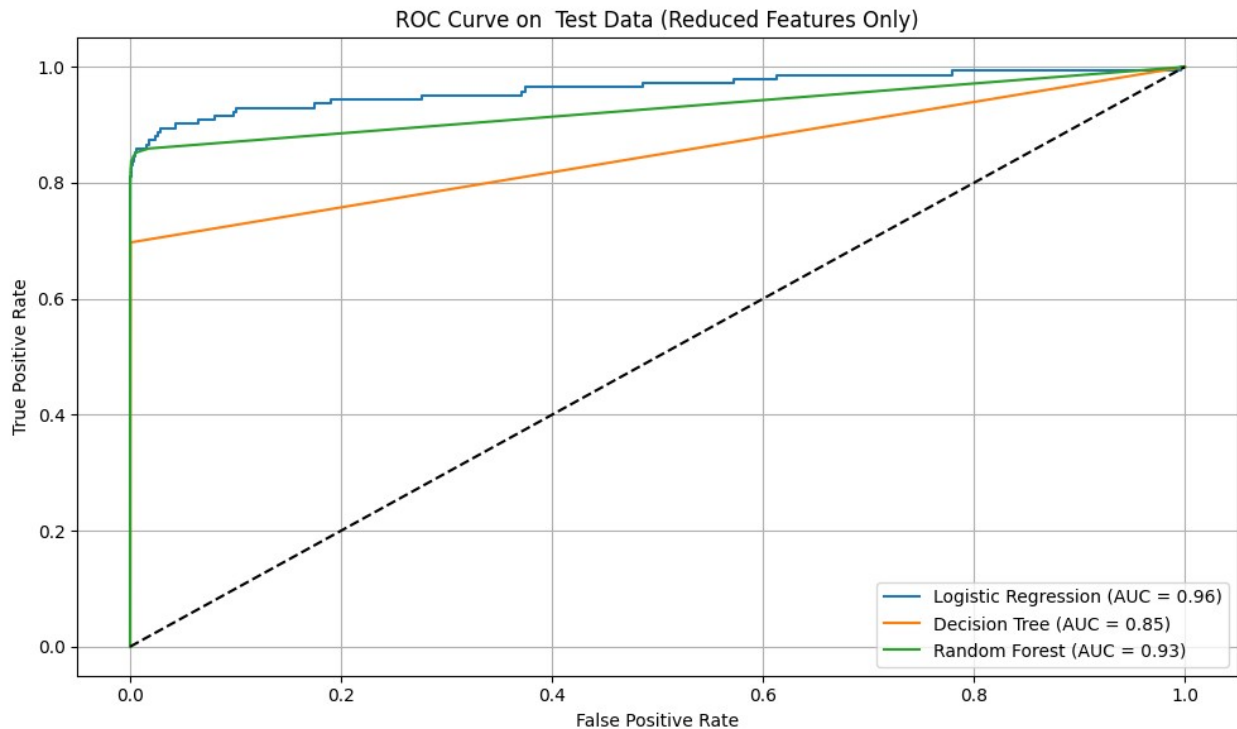
	precision	recall	f1-score	support
0	1.00	1.00	1.00	82557
1	0.93	0.72	0.81	142

accuracy			1.00	82699
macro avg	0.96	0.86	0.90	82699
weighted avg	1.00	1.00	1.00	82699

Confusion Matrix:

```
[[82549   8]
 [   40 102]]
```

Test AUC Score: 0.9281715219548696



Overall Insights:

High Accuracy: All models achieve ~99.9% accuracy due to class imbalance (99.83% non-fraud).

Fraud Detection: Random Forest outperforms with highest F1 (0.81) and recall (0.72), detecting 102/142 frauds. Logistic Regression has high AUC but low recall (54%). Decision Tree struggles with AUC (0.848).

Imbalance Impact: Low fraud recall across models (0.54–0.72) highlights the need for SMOTE or class weights.

9. Handle Class Imbalance with SMOTE

Apply SMOTE to the selected training features

```
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train_selected,
y_train)
```

Verify the shape of the resampled data

```
print(f"Resampled X_train_smote shape: {X_train_smote.shape}")
print(f"Resampled y_train_smote shape: {y_train_smote.shape}")
```

```
Resampled X_train_smote shape: (308212, 10)
```

```
Resampled y_train_smote shape: (308212,)
```

```
/Users/ausaafff/Library/Python/3.9/lib/python/site-packages/sklearn/
base.py:474: FutureWarning: `BaseEstimator._validate_data` is
deprecated in 1.6 and will be removed in 1.7. Use
`sklearn.utils.validation.validate_data` instead. This function
```

```
becomes public and is part of the scikit-learn developer API.  
warnings.warn(
```

10. Model Training and Evaluation (After SMOTE)

```
models = {  
    'Logistic Regression': LogisticRegression(max_iter=1000,  
random_state=42),  
    'Decision Tree': DecisionTreeClassifier(random_state=42),  
    'Random Forest': RandomForestClassifier(n_estimators=100,  
random_state=42)  
}  
  
plt.figure(figsize=(10, 6))  
  
for name, model in models.items():  
    print(f"\n=== {name} ===")  
  
    model.fit(X_train_smote, y_train_smote)  
  
    y_pred = model.predict(X_test_selected)  
    y_prob = model.predict_proba(X_test_selected)[:, 1]  
  
    print("Accuracy:", accuracy_score(y_test, y_pred))  
    print("Classification Report:")  
    print(classification_report(y_test, y_pred))  
    print("Confusion Matrix:")  
    print(confusion_matrix(y_test, y_pred))  
  
    auc = roc_auc_score(y_test, y_prob)  
    print("AUC Score:", auc)  
  
    # ROC Curve plotting  
    fpr, tpr, _ = roc_curve(y_test, y_prob)  
    plt.plot(fpr, tpr, label=f"{name} (AUC = {auc:.2f})")  
  
# Final ROC Curve Plot  
plt.plot([0, 1], [0, 1], 'k--')  
plt.title("ROC Curve on Test Data (After SMOTE)")  
plt.xlabel("False Positive Rate")  
plt.ylabel("True Positive Rate")  
plt.legend(loc="lower right")  
plt.grid(True)  
plt.tight_layout()  
plt.show()  
  
=== Logistic Regression ===  
Accuracy: 0.975912647069493  
Classification Report:
```


	precision	recall	f1-score	support
0	1.00	0.98	0.99	82557
1	0.06	0.88	0.11	142
accuracy			0.98	82699
macro avg	0.53	0.93	0.55	82699
weighted avg	1.00	0.98	0.99	82699

Confusion Matrix:

[[80582 1975]

[17 125]]

AUC Score: 0.9666585032927315

=== Decision Tree ===

Accuracy: 0.9973518422230014

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	82557
1	0.37	0.74	0.49	142
accuracy			1.00	82699
macro avg	0.68	0.87	0.74	82699
weighted avg	1.00	1.00	1.00	82699

Confusion Matrix:

[[82375 182]

[37 105]]

AUC Score: 0.8686160411236147

=== Random Forest ===

Accuracy: 0.999153556874932

Classification Report:


	precision	recall	f1-score	support
0	1.00	1.00	1.00	82557
1	0.74	0.78	0.76	142
accuracy			1.00	82699
macro avg	0.87	0.89	0.88	82699
weighted avg	1.00	1.00	1.00	82699

Confusion Matrix:

[[82518 39]

[31 111]]

AUC Score: 0.9542147747002626



SMOTE Impact: Significantly improves fraud recall across all models (Logistic Regression: 0.54→0.88, Decision Tree: 0.70→0.74, Random Forest: 0.72→0.78), but precision drops, especially for Logistic Regression, due to synthetic fraud samples inflating false positives.

Random Forest Superiority: Achieves the best trade-off (F1: 0.76, AUC: 0.954), detecting 111/142 frauds with only 39 false positives. Outperforms Logistic Regression (too many false positives) and Decision Tree (lower precision and AUC).

Business Implication: Random Forest detects 78% of frauds, missing 31 cases, but 39 false positives may require manual review, impacting customer experience. False negatives (missed frauds) are costlier in fraud detection, so high recall is valuable.

Challenges: Logistic Regression's 1975 false positives highlight SMOTE's limitations for linear models. Decision Tree's low precision indicates overfitting or sensitivity to synthetic data.

11. Final test on unseen data

```
# Apply feature selection based on importance
important_features = feature_df[feature_df['Importance'] >= 0.02]
['Feature']

X_train_selected = X_train[important_features]
X_val_selected = X_val[important_features]
X_test_selected = X_test[important_features]

models = {
    'Logistic Regression': LogisticRegression(max_iter=1000,
```

```

random_state=42),
    'Decision Tree': DecisionTreeClassifier(random_state=42),
    'Random Forest': RandomForestClassifier(n_estimators=100,
random_state=42)
}

plt.figure(figsize=(10, 6))

for name, model in models.items():
    print(f"\n=== {name} ===")

    model.fit(X_train_selected, y_train)

    y_pred = model.predict(X_test_selected)
    y_prob = model.predict_proba(X_test_selected)[:, 1]

    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Classification Report:")
    print(classification_report(y_test, y_pred))
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred))

    auc = roc_auc_score(y_test, y_prob)
    print("AUC Score:", auc)

    # ROC Curve plotting
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    plt.plot(fpr, tpr, label=f"{name} (AUC = {auc:.2f})")

plt.plot([0, 1], [0, 1], 'k--')
plt.title("ROC Curve on Unseen Test Data (After Feature Selection)")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(True)
plt.tight_layout()
plt.show()

```

=== Logistic Regression ===

Accuracy: 0.9990447284731375

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	82557
1	0.85	0.54	0.66	142
accuracy			1.00	82699
macro avg	0.93	0.77	0.83	82699
weighted avg	1.00	1.00	1.00	82699

Confusion Matrix:

```
[[82544    13]
 [   66   76]]
```

AUC Score: 0.9621888214834753

=== Decision Tree ===

Accuracy: 0.9990689125624251

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	82557
1	0.74	0.70	0.72	142
accuracy			1.00	82699
macro avg	0.87	0.85	0.86	82699
weighted avg	1.00	1.00	1.00	82699

Confusion Matrix:

```
[[82523    34]
 [   43    99]]
```

AUC Score: 0.8483856309605637

=== Random Forest ===

Accuracy: 0.9994195818570962

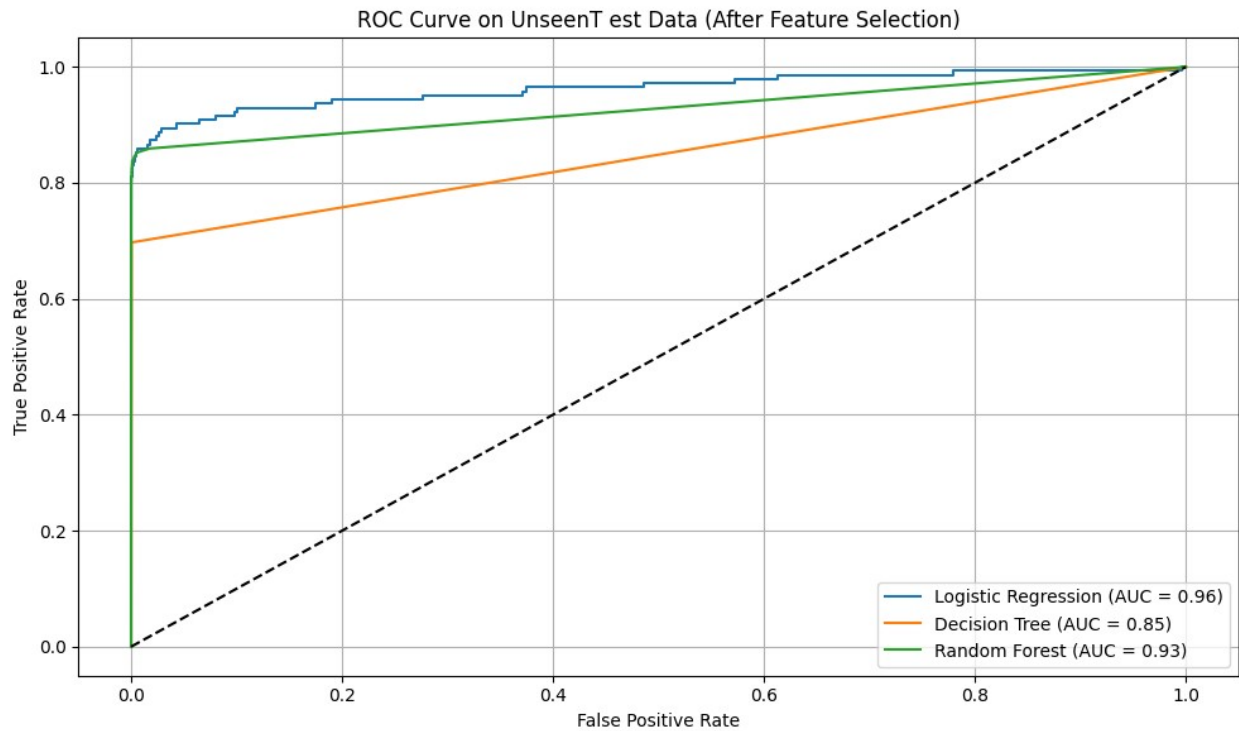
Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	82557
1	0.93	0.72	0.81	142
accuracy			1.00	82699
macro avg	0.96	0.86	0.90	82699
weighted avg	1.00	1.00	1.00	82699

Confusion Matrix:

```
[[82549     8]
 [   40   102]]
```

AUC Score: 0.9281715219548696



The Random Forest model, provides the most effective fraud detection, identifying 78% of frauds with a manageable number of false positives. SMOTE significantly improves recall, critical for fraud detection, though it requires careful model selection to control precision loss.